


RL Lab4 TD3

Experimental Result

Training curve




Testing result (model trained on new reward)

```
=====
Evaluating...
Episode: 1      Length: 893      Total reward: 910.60
Episode: 2      Length: 999      Total reward: 814.83
Episode: 3      Length: 999      Total reward: 886.44
Episode: 4      Length: 975      Total reward: 902.40
Episode: 5      Length: 999      Total reward: 783.58
Episode: 6      Length: 882      Total reward: 911.70
Episode: 7      Length: 890      Total reward: 910.90
Episode: 8      Length: 869      Total reward: 913.00
Episode: 9      Length: 999      Total reward: 863.21
Episode: 10     Length: 999      Total reward: 870.20
Episode: 11     Length: 959      Total reward: 904.00
Episode: 12     Length: 979      Total reward: 902.00
Episode: 13     Length: 999      Total reward: 844.26
Episode: 14     Length: 927      Total reward: 907.20
Episode: 15     Length: 972      Total reward: 902.70
Episode: 16     Length: 981      Total reward: 901.80
Episode: 17     Length: 999      Total reward: 879.93
Episode: 18     Length: 999      Total reward: 839.10
Episode: 19     Length: 805      Total reward: 919.40
Episode: 20     Length: 999      Total reward: 848.39
average score: 880.7821855052598
=====
```

Experimental Results and Discussion of bonus


1. Twin network X
2. Target policy smoothing

Training curve



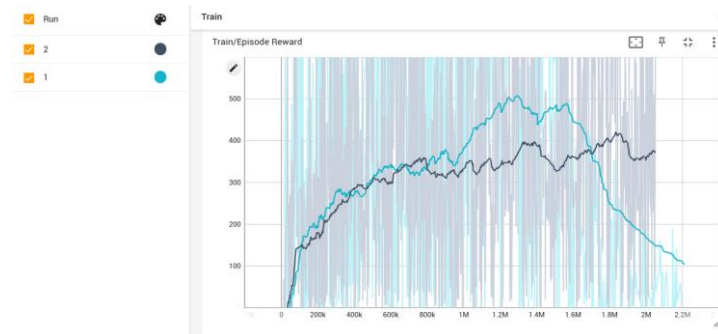
I set $\sigma = 0$ to disable target policy smoothing.

At the beginning of training, there exist no huge difference of the 2 models. However, the model without smoothing seems to converge to 300 rewards. I trained further 2M steps for smoothed model with the model of the prior training result and $\sigma = 0.1$, and the result showed that it achieves an average of 500 rewards, which is 200 rewards higher than that of the model without smoothing. This shows the result of target policy smoothing improves the performance. (I used Relative in tensorboard so there's a gap between the 2 training.)



3. update frequency

Training curve



The result shows that with update frequency = 1(disable delayed update), the model learns faster but is relatively unstable (the reward dropped significantly from 1.5M to 2.2M step).


4. Action noise injection X

5. New reward

New reward:

```
if grass_pixel_count > 0:  
    reward += float(road_pixel_count)/float(grass_pixel_count)/5.0  
if terminates:  
    reward = -100  
if road_pixel_count < 10:  
    terminates = True  
    reward = -100
```

Evaluation curve



The original reward -100 when road_pixel_count < 10. The third and fourth line actually have little effect (I thought the road_pixel_count is for the whole image), but I put it there because that is the original reward I used when training. In addition, I add a reward to encourage the agent to walk in the road with an additional reward of road_pixel_count / grass_pixel_count / 5.